

Weekly Briefing — Issue #1

Everything that mattered in AI this week, and what you should do about it.

Week of May 18 – 24, 2026 | Prepared for Jaroslaw

The 60-second version

- **The week in one idea: the industry pivoted to “the agent.”** At Google I/O and across the labs, the unit of work shifted from a chat turn to a supervised *fleet of agents* — model + runtime + a consumer agent, shipped together.
- **Google I/O 2026 (May 19) went all-in on agents** (~100 announcements): **Gemini 3.5 Flash** (frontier-ish at Flash price/speed), **Omni** (any-input→any-output), **Antigravity 2.0** (agent-first dev platform), Managed Agents in the API, and **Spark** — a 24/7 personal agent.
- **Anthropic answered with enterprise infrastructure:** self-hosted sandboxes (run agents inside your own VPC) and MCP Tunnels for private tool access — the data-residency story enterprises were blocked on.
- **DeepSeek made its 75% V4-Pro price cut permanent (May 24)** — frontier-class open-weight reasoning at ~¼ the cost of comparable US tokens; brutal downward pressure on inference pricing.
- **Learn this week:** writing **Skills** (markdown runbooks for agents), **context engineering** over prompt engineering, and **sub-agent + planning** orchestration — the durable leverage as models get stronger.

1 · The week in one idea

From “a better chatbot” to “a supervised fleet of agents”

The throughline of the week — anchored by Google I/O on May 19 — was that every major lab now ships an *agent stack*, not just a model: a frontier model, a **runtime** to run it autonomously (sandboxes, managed agents, tunnels), and increasingly a **consumer agent** (Spark). The job is moving from line-by-line prompting to **orchestration + verification**.

What to watch from here

- **Consumer agents go mainstream:** Gemini Spark is the most direct challenge to ChatGPT yet — expect a wider rollout within weeks.
- **Compute stays THE bottleneck:** the Blackstone–Google ~\$5B TPU JV (May 18) is one more sign capital + power gate the frontier.
- **Open-weight price pressure intensifies:** DeepSeek’s permanent cut resets build-vs-buy math for high-volume workloads.

2 · Model & Capability Landscape — what to use for what

I/O reframed the competition around **price/speed + distribution**, not just top-of-benchmark. The practical map this week:

When you need...	Reach for	Why
High-volume / latency-sensitive / cheap	Gemini 3.5 Flash	Frontier-ish quality at Flash price/speed (~■ the cost of comparable frontier models); powers Spark
Consumer “just do it for me”	Gemini Spark	Ambient 24/7 personal agent that reasons across your connected apps
Any-input → any-output multimodal	Gemini Omni	Generates across modalities (SynthID-watermarked output)
Cheapest tokens / self-host leaning	DeepSeek V4-Pro	75% cut now permanent — ~¼ the cost of equivalent US tokens

When you need...	Reach for	Why
Agent-first dev workflow	Google Antigravity 2.0 · Anthropic sandboxes	Multi-agent orchestration + VPC-local execution for sensitive data

Leaderboard scores vary by source — treat as directional. The week's story isn't a new #1; it's that the **runtime** around the model is now the battleground.

3 · Techniques & Skills to Learn (your leverage)

Study this section — it's where the compounding edge is as models get stronger.

“Skills” — the highest-leverage thing to learn now

[skill]

A **Skill** is a markdown file an agent loads for a specific task — an onboarding doc / runbook for an AI. Encode the procedure once; the agent follows it instead of you re-prompting. (Google also shipped **Science Skills** at I/O — the pattern is going mainstream.)

- **Write a runbook, not a question.** Anticipate failure modes, say which tool to use when, and specify **when to stop**.
- **Structure with sections + tags** (`## instructions`, `## tool_guidance`, `## output`) — models follow delineated structure far more reliably.
- **Include 2–5 canonical few-shot examples** of the exact behavior wanted.

Context engineering > prompt engineering

[practice]

Treat **context as a finite, precious resource**. Smarter models need *less* prescriptive prompting but *better* context curation — give the agent exactly what it needs, when it needs it, not everything.

Sub-agents + a forced planning phase

[practice]

Force a **planning phase before generation** (practitioners report complex-coding success jumping from ~■ to ~■ just by planning first), and use specialized **sub-agents** with single jobs — one session doing everything causes context pollution.

4 · Tools & Agentic Platforms

[tool]

- **Google Antigravity 2.0** — agent-first dev platform (desktop app, CLI, SDK, Managed Agents) with multi-agent orchestration.
- **Anthropic self-hosted sandboxes** (public beta) — run agents inside your own VPC (Cloudflare/Daytona/Modal/Vercel day one) + **MCP Tunnels** for private tool access.
- **MCP (Model Context Protocol)** — the standard way to connect agents to GitHub, Slack, Drive, etc. Wire up 1–2 MCP servers to your coding agent this week.
- **Gemini creative integrations** — Adobe (Firefly), Canva (Magic Layers), CapCut now inside Gemini.

Do this: Pick one repetitive weekly workflow and turn it into a Skill + sub-agent. Best ROI move right now.

5 · Market, Money & Business

Move	Detail	Signal
Compute land-grab	Blackstone–Google ~\$5B TPU joint venture (May 18)	Capital + power remain the gating moat
Open-weight pricing	DeepSeek V4-Pro 75% cut made permanent (May 24)	Downward pressure on inference cost is structural, not promotional
Distribution as moat	Google pushing agents to billions of devices via Gemini/Android	Reach, not just benchmarks, decides consumer-agent winners

Fast-moving figures are directional. The durable pattern: capability is becoming table stakes; **distribution + runtime + price** decide outcomes.

6 · This Week's Action List

- Convert one recurring workflow into a **Skill** (markdown runbook + few-shot examples).
- Add a **planning phase** to your agent prompts; split big jobs into **sub-agents**.
- Wire up **1–2 MCP servers** to your coding agent.
- Pilot **Gemini Spark** (or ChatGPT agent mode) on one real personal workflow.
- Re-run your inference cost math against DeepSeek V4-Pro's now-permanent price.
- For sensitive data, evaluate **VPC-local agent execution** (Anthropic sandboxes) over vendor-cloud.

Sources

blog.google (I/O 2026: Gemini 3.5 Flash, Omni, Spark, Antigravity, Science Skills) · CNBC · Anthropic (self-hosted sandboxes; MCP Tunnels) · DeepSeek · digitalapplied tracker · buildfastwithai daily roundups (May 18–24) · llm-stats.com (figures directional where noted).

Weekly Briefing — Issue #1 · Week of May 18 – 24, 2026 · Compiled from multiple cross-checked sources. Figures for fast-moving deals are directional.