

The AI Edge — Weekly Briefing

Everything that mattered in AI this week, and what you should do about it.

Week of May 24 – May 30, 2026 | Prepared for Jaroslaw | Issue #1

The 60-second version

- **Anthropic released Claude Opus 4.8** (May 28) with **Dynamic Workflows** — one agent can plan and run up to **1,000 parallel subagents**, do codebase-scale migrations, then self-verify. Fast mode is 3× cheaper.
- **Anthropic became the world's most valuable AI startup** — a **\$65B Series H** at **~\$965B**, passing OpenAI and nearing \$1T.
- **Google I/O 2026** (May 19) went all-in on agents: **Gemini 3.5 Flash**, plus **Spark** (a 24/7 personal agent) and **Antigravity**, an agent-first dev platform.
- **The industry pivoted to "the agent."** Three major labs changed their default model within a month. The battleground is autonomous, long-running, multi-agent work — not chat.
- **Learn this week:** context engineering, sub-agent orchestration, and writing **Skills** (markdown runbooks for agents). See Section 2.

1. Model Releases & Updates

Claude Opus 4.8 + Dynamic Workflows

Anthropic shipped Opus 4.8 on **May 28**, just 41 days after 4.7. Headline features:

- **Dynamic Workflows (research preview):** Claude plans a large task, spins up **hundreds of parallel subagents** (capped at 1,000) in one session, runs them longer, then **verifies its own outputs**. Demoed codebase-scale migrations across 100,000s of lines — kickoff to merge — using the test suite as the bar.
- **Effort controls:** a new control next to the model selector dials how much reasoning effort Claude spends.
- **Cheaper fast mode:** ~2.5× speed, now **3× cheaper** than on prior models. Pricing unchanged (~\$5 / \$25 per 1M tokens).
- GA in GitHub Copilot day one; positioned with better honesty / less deception.

Why it matters: This is the clearest sign yet that the unit of work is shifting from "a chat turn" to "a supervised fleet of agents." Your job becomes *orchestration* + *verification*, not line-by-line prompting.

Do this: In Claude Code, try a real migration/refactor with planning enabled and watch the decomposition into subagents. Use **low effort** for routine edits, **high** for architecture.

GPT-5.5 Instant becomes ChatGPT's default

On May 5, OpenAI made **GPT-5.5 Instant** the default across every ChatGPT tier, replacing GPT-5.3. Reported to cut hallucinated claims by 50%+ in some high-stakes scenarios.

Do this: Re-test saved prompts/automations against the new default — behavior and refusal boundaries shift on model swaps.

Google's Gemini: 3.5 Flash + Spark

At I/O (May 19) Google launched **Gemini 3.5 Flash** (beats 3.1 Pro on coding/agentive/multimodal, ~4× faster output), and **Spark** — a 24/7 personal agent that acts across your digital life (Gemini app, soon email/chat). 3.5 Pro arrives next month.

Why it matters: Google is using its distribution (billions of devices) to make agents ambient. Spark is the most direct consumer-agent challenge to ChatGPT yet.

Where the models stand (selected benchmarks)

Benchmark / area	Leader	Note
Scientific reasoning (GPQA Diamond)	Gemini 3.1 Pro — 94.3%	Top scientific reasoner
SWE-bench Pro (coding)	GLM-5.1 (Z.AI) — 58.4%	First open-weight model to hit #1 (Apr 7), since contested
Reasoning leaderboards	Claude Opus 4.8 (93), Qwen3.7 Max (92)	Best open-weight: DeepSeek V4 Pro (87)
Research / grounded synthesis	GPT-5.5 Pro	Retrieval + attribution + multi-doc synthesis

Leaderboard scores vary by source and methodology — treat as directional. "Claude Mythos Preview" tops some boards but is an unreleased preview.

2. New Techniques, Skills & Best Practices

This is the section to actually study — it's where your day-to-day leverage comes from.

"Skills" — the highest-leverage thing to learn now

A **Skill** is a markdown file an agent loads for a specific task — an **onboarding doc / runbook for an AI**. Encode the procedure once; the agent follows it instead of you re-prompting. (Google also shipped **Science Skills** at I/O — the pattern is going mainstream.)

How to write a good Skill:

- **Write a runbook, not a question.** Anticipate failure modes, define which tool to use when, and specify **when to stop**.
- **Structure with sections + tags:** XML/Markdown headers like `<background>`, `## instructions`, `## tool_guidance`, `## output`. Models follow delineated structure far more reliably.
- **Include 2–5 canonical few-shot examples** of the exact behavior wanted. Few-shot is still top-tier even with frontier models.
- **Be explicit about memory:** state what to remember across steps and what to discard.

Context engineering > prompt engineering

The 2026 consensus (Anthropic engineering): treat **context as a finite, precious resource**. Smarter models need *less* prescriptive prompting but *better* context curation. Don't dump everything in — give the agent exactly what it needs, when it needs it.

Sub-agents, planning & "dreaming"

- **Force a planning phase before generation.** Practitioners report success on complex coding tasks jumping from ~1/3 to ~2/3 just by planning first.
- **Use specialized sub-agents** with single, well-defined jobs. One session doing everything causes "context pollution" and worse output.
- **Anthropic's "dreaming" (May 6):** a scheduled, asynchronous memory-curation process — agents review past sessions, merge duplicate memories, drop stale ones, and surface recurring mistakes/preferences. Legal AI platform Harvey reported a **6×** improvement in task completion after enabling it. Points toward agents that genuinely learn your workflow.

The core prompting toolkit (keep these sharp)

- **Chain-of-Thought** — reason step-by-step before answering.
- **Self-Consistency** — sample multiple reasoning paths, take the majority; ~2–3× accuracy over plain CoT on hard problems.
- **ReAct** — interleave reasoning + tool actions; the backbone of most agents.
- **Few-shot + role prompting** — still reliable accuracy boosters.

3. Tools & Agentic Platforms

- **Google Antigravity** — agent-first dev platform (desktop app, agy CLI, SDK, Managed Agents). **Multi-agent orchestration** from day one: multiple agents on different parts of a codebase at once.
- **Anthropic Managed Agents** — added public-beta self-hosted sandboxes and a research-preview "MCP tunnels" feature (May 19).
- **MCP (Model Context Protocol)** is the standard way to connect agents to GitHub, Slack, Drive, Asana, etc. Wire up 1–2 MCP servers to your coding agent this week.
- **Claude Code vs Cursor**: Claude Code = terminal-native, autonomous multi-step + full-codebase work; Cursor = inline editing + autocomplete. Most shipping teams use **both**.

Do this: Pick one repetitive weekly workflow and turn it into a Skill + sub-agent. Best ROI move right now.

4. Market, Money & Business

Event	Details	Why it matters
Anthropic Series H	\$65B raised; ~\$965B valuation; passed OpenAI, nears \$1T (tripled from \$380B in Feb)	Capital consolidating around frontier-lab leaders
Anthropic × xAI compute	\$1.25B/month through May 2029; potentially \$40B+ to xAI	Compute is the real bottleneck — even rivals buy from each other
Cognition	Raised \$1B at \$26B valuation	Agentic dev tooling is a top funding magnet
Shield AI	\$1.5B Series G, \$12.7B valuation (+140% YoY)	Defense AI is surging
JPMorgan	~\$19.8B tech budget; AI reclassified R&D → core infrastructure	Enterprise AI: experiment → infrastructure
Novo Nordisk × OpenAI	AI across drug discovery, trials, manufacturing, supply chain	Vertical full-stack AI adoption in pharma
Apple (reported)	May let users pick 3rd-party providers (Google, Anthropic) for Apple Intelligence in OS 27	Apple conceding the model layer; distribution play

Big picture: Q1 2026 venture funding hit record highs (~\$300B+ globally), heavily AI-driven — flowing to frontier research, agent infrastructure, defense, and vertical tools for regulated industries. Some mega-deal figures (SpaceX/xAI structuring, a reported SpaceX IPO) are fast-moving and forward-looking — treat as developing.

5. Policy & Regulation

EU AI Act "omnibus" political agreement (May 7) — first major amendments since adoption:

- High-risk (Annex III) obligations **delayed** from Aug 2026 → **Dec 2027**.
- Two **new prohibitions**: AI generating non-consensual intimate imagery and CSAM.
- Synthetic-content marking delayed to Dec 2026; transparency rules still arrive **Aug 2026**.

US: the Dec 11, 2025 Executive Order aims to consolidate AI oversight federally and push back on state-law patchwork. State laws still landing: California's frontier-AI transparency act (Jan 1, 2026) and Colorado AI Act (Jun 30, 2026). A reported White House EO would create a **voluntary** framework for pre-release government model access.

Do this: EU-facing teams: the **Aug 2026** transparency rules (label AI-generated content) are the near-term deadline. US teams: track state-level laws, not just federal.

6. This Week's Action List

- Try Claude Opus 4.8 on a real multi-step task; learn the **effort control**.
- Convert one recurring workflow into a **Skill** (markdown runbook + few-shot examples).
- Add a **planning phase** to agent prompts; split big jobs into **sub-agents**.
- Wire up **1–2 MCP servers** to your coding agent.
- Re-test critical prompts against GPT-5.5 / Gemini Spark defaults.
- If EU-facing: plan for **Aug 2026** synthetic-content transparency labeling.

Sources

Anthropic (Opus 4.8; context-engineering; dreaming) · TechCrunch · VentureBeat · The New Stack · MarkTechPost · GitHub Changelog · blog.google (I/O 2026: 100 announcements, Spark, Antigravity, Science Skills) · 9to5Google / Tom's Guide · Tech Startups (Anthropic \$965B) · Crunchbase News (funding rounds) · SD Times · MarketingProfs (GPT-5.5, JPMorgan, Novo Nordisk, Apple) · YourStory / MindStudio (dreaming; Harvey 6x) · llm-stats.com / Vellum / BenchLM (benchmarks) · Global Policy Watch & White & Case (EU AI Act + US EO) · f22labs / Totalum (Claude Code productivity).

The AI Edge · Issue #1 · Week of May 24–30, 2026 · Compiled from multiple cross-checked sources. Figures for fast-moving deals are directional, not guaranteed.