

Weekly Briefing — Issue #5

Everything that mattered in AI this week, and what you should do about it.

Week of Jun 22 – Jun 28, 2026 | Prepared for Jaroslaw

The 60-second version

- **The week in one idea: the work became delegable — so your edge moved to the guardrail.** Last week's move was *subtract* the middle layer; this week you got to hand the actual *doing* to a machine — a design agent that builds to your system, a browser agent that clicks through your dashboards, a coding agent that takes real actions, a bid strategy that chases revenue, an analysis that names your biggest leak. The catch is identical every time: each delegation only ships safely because **you set one control** — a Skill the agent obeys, a sandbox, a human-approval gate, a campaign type you can still see into, a fallback model. The operator job stopped being 'do the task' and became 'write the rule and hold the gate.'
- **Design + Dev — Config 2026 and the agent stack made 'declare it once, let the system execute' the week's pattern.** Figma's on-canvas **design agent** went open-beta with **Skills** — write your design system once in Markdown (/follow-ds-guidelines) and every teammate's agent designs to your tokens, not generic defaults; **Figma Motion** animates on the canvas and copies the exact easing out as CSS/React; **Shader fills** turn 'find a grain texture' into 'describe the grain.' On dev: **Vercel WebSockets** (realtime with no separate backend), **Next.js 16.3 Instant Navigations** (SPA-feel from two flags), and **AI SDK 7** (agents that pause for a human yes/no and survive a restart). The thread: stop translating between tools and stop re-explaining yourself — state the intent once.
- **AI tooling — autonomy shipped, but every drop came with a leash.** Google moved **computer use** into **Gemini 3.5 Flash** (\$1.50/\$9 per M tokens, ~■ of GPT-5.5) so a cheap model can click through the GUI-bound grunt work that has no API — but only behind a Docker sandbox, an allowlist, and injection auto-halt. Claude Code's **/loop** and **/goal** let an agent run unattended — but **/goal** only stops when a separate cheap evaluator confirms your condition. The pattern under both: hand over the doing, keep a gate that catches a silent failure.
- **Marketing — let the system optimize; your edge is the input and the real signal.** Google Search Console now shows your **first-party** impressions inside AI Overviews/AI Mode — stop guessing your AI-search presence with third-party estimators and baseline the source of truth. **Standard Shopping** can finally bid for revenue (Maximize conversion value, no required tROAS) without surrendering to the PMax black box. **Meta multimedia ads** want one ad packed with up to 10 mixed assets so delivery picks the winner. And the cheapest email win is *shrinking* the list: a behaviour-based **re-engagement** → **sunset** flow protects inbox placement for your buyers.
- **News — the frontier became a government decision, and compute got rationed.** Anthropic's **Mythos 5** came back but *restricted* to ~100+ US critical-infrastructure operators (Fable 5 general access still dark); OpenAI's **GPT-5.6** (Sol/Terra/Luna) ships via a **customer-by-customer** government-approved preview; Google **capped Meta's use of Gemini** because it couldn't supply the capacity. Open weights filled the vacuum — Z.ai's **GLM-5.2** drew 'matches the US frontier' headlines (read with the harness-and-scaffolding asterisk). The durable read: the best model may not be in your account on day one — design around the tier you can actually get, and keep an open-weight fallback you can run yourself.
- **Sharpen your edge:** the connective thread to last week ('subtract the middle layer') is direct — once a layer is gone, the work it used to gate gets handed to an agent or a default, and the only thing left for you to own is the guardrail. Every play below names that guardrail explicitly: a Skill, a sandbox, a human-approval gate, a spend cap, a fallback model, a re-measure date. Take the automation. Write the rule. Hold the gate.

The week in one idea: hand off the doing, own the guardrail

The wins this week were delegations, not just subtractions

Line the week's plays up and they rhyme. Figma's design agent *builds the screen* — you just hand it a Skill (your design system as a Markdown slash command) so it builds to your tokens. Gemini 3.5 Flash *clicks through the dashboard* that has no export button. AI SDK 7 agents *charge the card and send the email*. Google's Standard Shopping *decides the bids* to chase revenue. AI *reads your order export* and names the single biggest dollar leak. Last week the cheapest win was removing a dependency, a server, a manual step; this week it was removing *yourself from the loop* — letting an agent or a platform default take the action it used to need you for. That's the escalation: from subtracting the middle layer to delegating the judgment that sat on top of it.

Synthesis, not recap: six-plus dailies this week — across design, dev, AI tooling, paid, growth and CRO — told one story from different crafts. The work itself became delegable.

But every delegation shipped with a control you set — and the frontier got gated from above

The discipline is the twin of the delegation, exactly like last week's 'every subtraction shipped with a guard.' Let the design agent build — but constrain it with a Skill. Let the browser agent click — but only in a Docker sandbox, behind an allowlist, with injection auto-halt on. Let the coding agent act — but gate the irreversible tools behind a human yes/no and make it durable so the human can say yes *hours* later. Let `/goal` run unattended — but only until a separate, cheap evaluator confirms the condition. Let Google bid for revenue — but in Standard Shopping, where you keep the search-term visibility a PMax black box takes away. And the news drew the same line from the top: the best models now arrive **gated** (Mythos restricted, GPT-5.6 approved customer-by-customer) and compute is **rationed** (Google capping Meta), so 'which model, is it cleared, what's the fallback' became a planning input, not a footnote.

Do this: This week's homework, one line each: write your design system as a Figma Skill and make one agent build to it (design); gate one money-spending tool in an AI SDK 7 agent behind human-approval before you ship it (dev); run one Gemini computer-use task in the Docker sandbox with the two opt-in safeguards ON (AI); open Search Console → Generative AI and date-stamp your AI-search baseline (growth); switch one Standard Shopping campaign to Maximize conversion value with the tROAS box empty (paid); and add one open-weight model (GLM-class) to your fallback list and test it on a real task now, not in a crisis (all).

What changed, and what to do — by domain

One tight synthesis per domain that moved this week. The throughline (delegate the doing, own the guardrail) shows up differently in each.

Web design — teach the agent your system once, then let it design and animate

[web-design · figma]

Config 2026 (Jun 24) made the Figma **design agent** open beta for everyone and gave it **custom Skills** — a reusable Markdown instruction set you invoke as a slash command, so a Skill named `follow-ds-guidelines` becomes `/follow-ds-guidelines` and the agent designs to *your* component library, token names, spacing scale and a11y minimums instead of generic defaults. That's the governance fix for the #1 failure of AI-in-design: output that ignores your system and forces rework. Same event shipped **Figma Motion** (an on-canvas animation timeline whose Dev Mode tab copies the exact easing as CSS/JSON/React — no After-Effects-to-code translation).

Do this: Write your design-system rules once as a Skill (Skills → Add skill → numbered instructions + one worked example) and test with `/follow-ds-guidelines` design a settings page. For motion, keyframe one button-hover in Figma Motion and copy the CSS from Dev Mode. Both need a paid plan + Full seat; both are free of AI credits during beta (credits apply at GA, so cost it before standardizing).

Keeper principle: a Skill turns 'AI that guesses' into 'AI that follows our system' — the cheapest governance lever in design right now, no engineering required. Distinct from Code Layers (closed beta, July — not run-it-today).

Web dev — the platform absorbed the backend, the rewrite, and the agent state machine

[web-dev · vercel · nextjs · ai-sdk]

Three independent dev drops, one root: let the platform hold the hard part. **Vercel WebSockets** (public beta) let a Function hold a two-way socket open with standard libraries and idle-cheap Active-CPU billing — realtime (chat, presence, token streaming) with no separate always-on box. **Next.js 16.3 Instant Navigations** turn two config flags (`cacheComponents`, `partialPrefetching`) into SPA-feel navigation without a SPA rewrite, with Instant Insights flagging slow routes as dev-time errors. **AI SDK 7** made agents production-safe: `toolApproval` pauses a risky tool

for a human yes/no (optionally HMAC-signed), and workflowAgent persists state so the agent survives a deploy, a crash, or an approval that lands hours later.

Do this: Add realtime to an existing Vercel project with one api/ws.ts (move shared state to Redis — instances aren't sticky). Prototype Instant Navigations on a branch (npm i next@preview) and fix the one route Instant Insights flags. For agents that spend/send/delete, upgrade to ai@7 (Node 22+/ESM), wrap actions in ToolLoopAgent, and gate the dangerous tools behind 'user-approval' — then pair a spend cap so an approved-but-looping agent can't blow the budget.

All three are vendor-confirmed (single-vendor primary), and two are pre-production previews (Next.js @preview, Vercel/Next WebSockets experimental) — prototype, don't ship to users yet. Treat your AI coding toolchain as attack surface; gate, sandbox, cap.

Graphic — describe the material instead of hand-painting it, and keep it live

[graphic · figma · webgpu]

Config 2026 also shipped Figma **Shader fills and effects** — real WebGPU shaders you create by describing the look to the Figma agent (or dropping a reference image): frosted glass, chrome, dithering, halftone, particles, plus stackable looks like Gradient map and Riso print. The win is that each shader is **parameterized by default**, so you get sliders on the canvas — a brand treatment becomes a reusable, tweakable component, not a baked PNG you can't adjust later.

Do this: For your next surface treatment (glassmorphism panel, halftone hero, film grain), stop hunting stock textures: select the layer, add a Shader fill (new material) or Shader effect (transform existing art), prompt the look, then dial the on-canvas sliders to your brand palette. Verify the look survives at export/handoff sizes. Paid plan + Full seat; free of credits in beta, charged at GA.

The graphic version of the week's theme: hand the tool structured intent and get back something live and editable, not a frozen pixel you re-make next quarter.

AI tooling — autonomy with a leash: a cheap browser agent and unattended loops

[ai-tooling · gemini · claude-code · agents]

Two run-it-today plays both gave you more autonomy and a required guard. Google moved **computer use into mainline Gemini 3.5 Flash** — one fast, cheap model (\$1.50 in / \$9 out per M tokens, 78.4 on OSWorld-Verified, a hair behind GPT-5.5) that sees a screen and clicks/types/scrolls — killing GUI-bound grunt work with no API (pulling numbers from dashboards that won't export, scheduled competitor screenshots, UI smoke-tests). And Claude Code's **/loop** (re-run on a schedule) + **/goal** (work until done) turned babysitting into a walk-away task — with **/goal** judged by a separate, cheap evaluator (defaults to Haiku).

Do this: Feel the Gemini loop with the zero-code Browserbase demo, then clone Google's Docker-sandboxed reference repo and — before any real task — flip ON the two opt-in safeguards (confirm-before-sensitive-actions + auto-halt-on-injection, both OFF by default) plus a website allowlist. For Claude Code, try **/goal all tests in test/auth pass and lint is clean or stop after 20 turns**, or drop a **.claude/loop.md** with your PR-watch instructions.

The binding caveat on computer-use is indirect prompt injection (malicious instructions hidden in a page) — Google says 'no single safeguard is sufficient,' so layer sandbox + allowlist + human-in-the-loop + injection auto-halt. **/loop** and **/goal** are session-scoped and expire after 7 days; graduate to Routines for machine-off scheduling.

Growth — measure your AI-search presence from the source of truth, not an estimate

[growth · aeo · search-console]

Google Search Console now exposes a **Generative AI performance report** — your impressions inside AI Overviews, AI Mode and Discover, segmentable by page/country/device — and Mueller clarified (Jun 25) how impressions are counted (a real link must be shown; a favicon doesn't count; the same site twice in one answer is one impression). This is the first-party replacement for third-party AEO estimators. It also resolved our Jun 21 call (Google ships a first-party AI-visibility metric ■) — with the honest caveat that the report actually launched Jun 3 and is impressions-only.

Do this: Open Performance → Generative AI, export your top and bottom URLs by AI-feature impressions, and date-stamp it as your baseline. Double down on the formats already winning citations (intent title → one-sentence answer → list/table → an original first-party stat → FAQ schema + last-updated) and rewrite the gaps to match. Re-measure on a weekly export so you can prove cause and effect.

Measure → fix → re-measure now runs on Google's own data, not a leap of faith. Read it with Mueller's counting rules so you don't over- or under-read a number.

Social + Paid — set the input and the guardrail; let the platform's AI optimize

[social · paid · meta-ads · google-ads]

The marketing plays all moved the lever to your inputs while the platform automated the optimization. Social: Meta's free **Creator Assistant** answers questions against your *own* Facebook account data ('why did this Reel win?') and brainstorms off live trends — a coach that sees your real numbers beats a blank prompt. Paid: Meta's refreshed **multimedia-ads** guidance says pack up to 10 mixed image+video assets (multiple aspect ratios) into *one* ad and let Advantage+ pick the winner per placement — so the edge is creative variety, not audience micro-targeting; and Google's **Standard Shopping** can now use Maximize conversion value *without* a required Target ROAS, so you bid for revenue while keeping the search-term visibility a feed-only PMax campaign takes away.

Do this: Social: ask Creator Assistant why your top Reel outperformed, extract the repeatable variable, then ask for 3 angles off what's trending (US/CA/IN only for now). Paid (Meta): build one ad with ~5 genuinely different creatives (UGC, demo, testimonial, explainer, lifestyle) across 9:16 + 1:1/4:5, Advantage+ creative ON, AI assets disclosed. Paid (Google): switch one Standard Shopping campaign to Maximize conversion value, leave tROAS empty for 2–3 weeks, then add a tROAS near your *achieved* ROAS to hold efficiency.

Same move as the growth story: the platform runs the optimization; your job is to set the inputs and the guardrails so its automation lands on your terms. Note the 'Maximize conversion value with a target ROAS' strategy is being renamed simply 'Target ROAS' — same behaviour, update your SOPs the same week.

Email — protect the inbox by pruning the dead list and closing the spoof hole

[email · deliverability · dmarc]

Two durable deliverability moves, both about *wantedness*. Litmus (Jun 2026): mailbox providers grade you on whether people want your mail, so a big pool of never-openers teaches Gmail/Yahoo/Outlook your sends are unwanted — and routes them to spam for your *buyers* too. The fix is a standing **re-engagement** → **sunset** flow keyed on behaviour (no click/visit/purchase in 6+ months, not opens — Apple Mail auto-opens make open-rate a lie). Separately, DMARC is now a full standard (**RFC 9989 / DMARCbis**): its np=reject tag rejects mail from subdomains that don't exist in your DNS — the phantoms spoofers ride to borrow your reputation.

Do this: Build the disengaged segment, run a 2–3 email win-back, then **suppress** the non-responders and prune hard-bounced/invalid addresses (some are spam traps). Measure the lift in Gmail Postmaster's 'Deliverability analysis' before/after. Separately, a 10-minute DNS edit: append np=reject to your _dmarc record (with p=quarantine/reject already in place) and delete deprecated pct/rf/r/i.

It feels scary to shrink a list on purpose, but a smaller awake list lands in the inbox and a big ignored one doesn't. np=reject is prerequisite-grade hygiene before any BIMi logo play.

CRO — let your own data name the single highest-dollar lever

[cro · aov · analytics]

A teammate's Ask-anything question ('the best ONE thing for CRO and AOV?') got the honest answer: it's whatever *your* numbers say is leaking most, and you can decide it in under an hour. Hand the AI your last-90-day order export + GA4 funnel and have it compute your AOV and your **modal** (most common) order value, then name the biggest leak. If it's basket size → a free-shipping threshold set just above your modal order value plus a 'You're \$X from free shipping' progress bar (the goal-gradient effect). If it's checkout drop-off → kill friction first (guest checkout, fewer fields, full cost shown early — surprise costs are the #1 documented abandonment reason, per Baymard).

Do this: Export the data, ask Claude/ChatGPT to compute AOV + modal order value and rank the leak, then ship *one* lever — peg the threshold to the modal value (not a round number, not the blunt '~30% above AOV' rule) — and measure AOV, conversion AND margin for 3–4 weeks before the next change.

Stop importing someone else's case study; let your distribution pick the lever. A higher AOV you fund with absorbed shipping can be a net loss — watch margin, not just AOV.

News — the frontier became a government decision; keep a fallback you can run yourself

[news · ai-tooling · markets · policy]

The week's biggest landscape shift: the 'newest model, instantly available to everyone' era ended for the frontier tier. Anthropic's **Mythos 5** was restored but *restricted* to ~100+ US critical-infrastructure operators (Fable 5 general access still suspended, no timeline) — resolving our Jun 13 restoration call ■ as a partial/restricted one. OpenAI's

GPT-5.6 (Sol/Terra/Luna) ships via a **customer-by-customer** government-approved preview. And capacity, not just capability, became a planning input: Google **capped Meta's use of Gemini** because it couldn't supply the compute. Into that vacuum, open weights advanced — Z.ai's **GLM-5.2** drew 'matches the US frontier at finding bugs' headlines, but the parity claim rests on a narrow Semgrep test and practitioners pinned it on scaffolding, not the base model — directional, not verified. On the markets side the month's repricing whipsawed (Jun 22–23 rout on capex/IPO fears → Jun 25 Micron blowout + Morgan Stanley 'crowded dip, not broken ROI', which resolved our Jun 23 capex-ROI-defence call ■).

Do this: Treat model availability as a procurement question: scope client work around the generally-available tier (Terra/Luna-, Sonnet-class), keep at least one open-weight model (GLM-class) in your fallback list and test it on real tasks now, and design features to degrade gracefully under rate limits and price drift. Trust the benchmark you ran, not the headline.

Climate signal, not a trade. The open 'dip not regime change' call (due ~Jul 20) stays open, leaning held but volatile; the Google→Meta cap supports the 'compute is the binding constraint' call (due ~Aug 3).

Techniques & tools worth adopting this week — the keepers

The survivors after killing 95%. Each is a run-it-today play with a real link and a built-in guardrail.

The shortlist

[\[technique\]](#)

Keeper	What it buys you	Start here + the guardrail
Figma agent Skills	The design agent builds to YOUR design system, not generic defaults	Skills → Add skill → DS rules as Markdown; test with /follow-ds-guidelines
Figma Motion → Dev Mode	Animate on-canvas and copy exact CSS/React easing (no After Effects)	Keyframe one hover; copy the CSS; validate the curve on your front-end
Gemini 3.5 Flash computer use	A cheap agent clicks through GUI work that has no API/export	Docker sandbox + allowlist + BOTH opt-in safeguards ON before any real task
AI SDK 7 tool approvals	Agents that pause for a human yes/no and survive a restart	toolApproval:'user-approval' on spend/send/delete tools + a spend cap
Next.js 16.3 Instant Navigations	SPA-feel navigation from two flags, no SPA rewrite	On a branch: cacheComponents + partialPrefetching; fix the route Instant Insights flags
GSC Generative AI report	First-party impressions inside AI Overviews/AI Mode	Performance → Generative AI; export + date-stamp a baseline; re-measure weekly
Standard Shopping max-value bidding	Bid for revenue without the PMax black box	Switch strategy, leave tROAS empty 2–3 wks, then set tROAS near achieved ROAS
Email sunset flow + DMARCBis np=reject	Protect inbox placement for your buyers; close the subdomain spoof hole	Behaviour-based win-back→suppress; append np=reject to _dmarc (10-min DNS edit)

Predictions ledger — resolved & newly on the board

Resolved this week

- **[news] Anthropic restores access to a suspended model** (made Jun 13) — the US government cleared **Mythos 5** for redeployment to ~100+ US critical-infrastructure operators (Commerce Sec. Lutnick: 'appropriate safeguards are in place'). Honest caveat: a *restricted* re-release — **Fable 5 general access stays suspended** with no timeline (new open call logged, due ~Sep 30).
- **[growth] Google ships a first-party AI-visibility metric in Search Console** (made Jun 21) — the Generative AI performance report (impressions inside AI Overviews/AI Mode/Discover; Mueller clarified counting Jun 25). Honest caveat: it had already launched Jun 3 (we missed it), and it's impressions-only, not a Bing-style

citation-*share* metric — early/partial on specifics, true on substance.

- ■ **[news]** A US AI-megacap or major sell-side note defends AI-capex ROI within ~2 weeks of the Jun 22–23 selloff (made Jun 23) — Morgan Stanley (Jun 24) reframed the drop as 'not expensive, but crowded' with the 'earnings-revision story validated'; Micron's blowout quarter (HBM sold out through 2027) was the supporting evidence.

Newly on the board (logged this week, still open)

- **[web-design]** Figma agent Skills become billable AI-credit actions at GA, AND/OR a community Skills marketplace / cross-tool .md skill standard emerges — due ~end of Q1 2027.
- **[web-dev]** A 2nd major JS agent framework (LangGraph, Mastra, OpenAI Agents SDK, LlamaIndex) ships BOTH built-in tool-approval and state-persisting durable execution, matching AI SDK 7 — due ~end of Q1 2027.
- **[web-dev]** Vercel WebSocket support for Functions exits beta to GA and Next.js drops the `experimental_` prefix — due ~end of Q3 2026.
- **[ai-tooling]** A 2nd frontier lab (OpenAI or Anthropic) ships a fast-tier computer-use/browser-control model at ≤ ~\$2/\$10 per M tokens to match Gemini 3.5 Flash — due ~end of Q3 2026.
- **[growth]** Google adds clicks and/or CTR (possibly query data) to the Search Console Generative AI report, impressions-only today — due ~end of Q4 2026.
- **[paid]** Google completes the 'Maximize conversion value + tROAS' → 'Target ROAS' rename platform-wide, AND/OR extends no-required-tROAS value bidding beyond Standard Shopping — due ~end of Q4 2026.
- **[email]** A major ESP ships an automated engagement-based suppression / 'sunset' default tied to a deliverability-risk signal — due ~end of Q1 2027.
- **[news]** Fable 5 general (unrestricted) access is restored (due ~Sep 30), and GPT-5.6 reaches GA without per-customer US approval (due ~Aug 15).
- **[news]** A 2nd hyperscaler/model vendor publicly rations or caps a major customer's AI compute within ~8 weeks, confirming compute — not capability — as the binding constraint — due ~Aug 23.

This week's action list

Concrete and assignable, per discipline. Each pairs the delegation with the guardrail.

- Design — Write your design system as a Figma Skill and make one agent build to it (/follow-ds-guidelines); copy one hover's easing from Figma Motion → Dev Mode.
- Dev — Gate one money-spending tool in an AI SDK 7 agent behind 'user-approval', make it durable, and put a spend cap in front of it.
- Dev — Prototype Next.js 16.3 Instant Navigations on a branch and fix the one route Instant Insights flags; add realtime via one Vercel `api/ws.ts` (state in Redis).
- AI/all — Run one Gemini 3.5 Flash computer-use task in the Docker sandbox with the two opt-in safeguards ON; wire one /goal loop with a clear 'done when' condition.
- AI/all — Add one open-weight model (GLM-class) to your fallback list and test it on a real task today, not in a crisis.
- Growth — Open Search Console → Generative AI, export top/bottom URLs by AI-feature impressions, and date-stamp your baseline.
- Paid — Switch one Standard Shopping campaign to Maximize conversion value (tROAS empty 2–3 weeks); pack one Meta ad with ~5 distinct creatives across 9:16 + 1:1/4:5.
- Email — Build the behaviour-based disengaged segment, run win-back → suppress, and append `np=reject` to your `_dmarc` record.
- CRO — Hand AI your 90-day order export + GA4 funnel; compute AOV + modal order value; ship ONE lever and measure AOV, conversion AND margin.
- Social — Ask Facebook Creator Assistant why your top Reel won, extract the variable, and turn 3 trend angles into the week's posts.

Sources

[Figma Blog — The Figma agent: custom tools, context, and skills](#) · [Figma Learn — Custom skills for the Figma agent](#) · [Figma Blog — Introducing Figma Motion](#) · [Figma Blog — Config 2026 recap \(shaders\)](#) · [Vercel Changelog — WebSocket support \(public beta\)](#) · [Next.js Blog — 16.3: Instant Navigations](#) · [Vercel Changelog — AI SDK 7](#) · [Vercel Blog — A new programming model for durable execution](#) · [Google — Computer use in Gemini 3.5 Flash](#) · [Google AI for Developers — Computer use \(Gemini API docs\)](#) · [Claude Code Docs — Run prompts on a schedule \(/loop\)](#) · [Claude Code Docs — Keep Claude working toward a goal \(/goal\)](#) · [Google Search Console Help — Generative AI performance report](#) · [Search Engine Journal — Mueller on counting AI-search impressions](#) · [Search Engine Land — Max Conversion Value bidding for Standard Shopping](#) · [Meta for Business — About multi-media ads](#) · [Meta Newsroom — Creator Assistant](#) · [Litmus — Re-engagement & list hygiene](#) · [IETF — RFC 9989 \(DMARCBis\)](#) · [Shopify — Average Order Value guide](#) · [Baymard — Cart Abandonment Rate](#) · [Engadget — US lets Anthropic redeploy Mythos](#) · [The Decoder — GPT-5.6 customer-by-customer approval](#) · [Bloomberg — Google caps Meta's use of Gemini \(FT\)](#) · [VentureBeat — Z.ai's open-weights GLM-5.2](#) · [CNBC — Micron blowout quarter \(HBM sold out through 2027\)](#)

Weekly Briefing — Issue #5 · Week of Jun 22 – Jun 28, 2026 · Compiled from multiple cross-checked sources. Figures for fast-moving deals are directional.