

Weekly Briefing — Issue #7

Everything that mattered in AI this week, and what you should do about it.

Week of Jul 6 – 12, 2026 | Prepared for Jaroslaw

The 60-second version

- **The week in one idea: the vendors turned the gauge into a pitch — 'Opus-class, much faster,' a voice that answers from the live web, a border property that redraws the box — so your edge is the ONE cheap test that proves each claim on your own ground before you trust it.** Last week the platforms cracked open and handed you a gauge to READ (SNDS, PMax Channel Diagnostics, your own eval on Sonnet 5); this week they handed you a louder *claim*, and the move is the same discipline one notch upstream: don't reprice, redraw or re-platform on a launch benchmark or a founder's tweet. Grok 4.5 is 'roughly Opus at a third of the price' — so route it behind your gateway and wait for the independent leaderboard, don't hard-code it. GPT-Live finally answers with current facts — so verify the fact it fetched. border-shape is 'here' — so read the deploy tier (it's flag-only) and ship the graceful-degrade half. And CRO's 'one thing' is whatever YOUR order data says leaks most. Take the upgrade; run the one test that earns your trust.
- **Web design — non-rectangular UI with no SVG or clip-path, but ship the half that degrades and guard the half that doesn't.** Two CSS properties kill the 'anything but a rounded rectangle costs an asset or a hack' tax. **corner-shape** (Chrome/Edge/Opera 139+) turns any border-radius into a **bevel, notch, scoop, squircle** or a `superellipse()` curve — up to four values for per-corner ticket/tab silhouettes — with **zero downside**: border-radius is the fallback, so unsupported browsers just show your rounded corner (no @supports needed). Ship it today. The bigger, near-future property is **border-shape**: it redraws the box geometry itself so **border, box-shadow, outline and pointer events all follow a fully custom edge** (speech-bubble tails, wavy banners) — but it's **flag-only in Chrome Canary 146+**, in no stable browser, syntax may still change, so wrap it in `@supports (border-shape: circle())` as pure enhancement. The gauge is the deploy tier: graceful-degrade → ship now; flag-only → guard it. Card: card-webdesign-border-shape.
- **CRO/AOV — the honest answer to 'the best ONE thing' is: there isn't one, so stop copying tactics and let AI read YOUR own numbers.** A standing team Ask got its run-it-today method: export your last-90-day order totals + your GA4 checkout funnel, hand both to Claude/ChatGPT, and ask it to compute your AOV and your **modal** (most-common) order value and say whether you lose more to **small baskets** or **checkout drop-off**. If it's drop-off (CRO), attack cost-shock — ~70% of carts abandon and the top fixable reasons are **unexpected extra costs (39%), forced account creation (19%)** and a **too-long checkout (18%)** (Baymard, 50-study dataset): show landed cost early, offer guest checkout, cut fields. If it's small baskets (AOV), set a free-shipping threshold just above your modal order value with a live 'You're \$X away' progress bar (goal-gradient). Ship ONE lever and measure AOV, conversion AND margin for 3–4 weeks — a higher AOV you fund with absorbed shipping can be a net loss. Card: card-cro-ai-diagnose-one-lever.
- **AI tooling — voice finally got a brain, but keep the human verify gate on what it fetches.** OpenAI shipped **GPT-Live** (Jul 8), a full-duplex voice partner that **listens while it talks** (interrupt it like a colleague; pause and it waits), backchannels ('mhm'), and for anything needing fresh facts **silently hands off to GPT-5.5 in the background** and keeps talking while the answer comes back — so voice finally has current knowledge + live search, fixing the old ~2024-cutoff voice mode that made real thinking useless. Run it today: update the ChatGPT app (or chatgpt.com), tap Voice (Free = GPT-Live-1 mini, any paid tier = GPT-Live-1), prime it once ('be my brainstorm partner; look things up and keep talking'), and think out loud on a brief while it pulls live references. Honest scope: **no screen-share/video/vision at launch, no API yet**, and it's **excluded from Business/Enterprise/Edu** — so keep confidential client material off it and verify the facts it fetches. Card: card-ai-tooling-gpt-live.

- **News — the model price war opened a new front and OpenAI's product org lost its owner, so 'which model' is now a routing decision, not a hard-code.** xAI shipped **Grok 4.5** (Jul 8), an 'Opus-class' coding model Musk pegs as 'roughly comparable to Opus 4.7 but much faster' at **\$2/\$6** per M tokens vs Opus 4.7's \$5/\$25 — the first major model since xAI's **\$60B all-stock Cursor buy**. A day later **Meta entered the coding wars with Muse Spark 1.1**, a low-cost agentic model on a new OpenAI/Anthropic-compatible API (Zuckerberg's first X post in ~3 years). **SambaNova** closed a **\$1B first tranche at an \$11B valuation** (Series F, General Atlantic) — capital still flooding inference silicon that isn't Nvidia. And **Fidji Simo** is stepping down from OpenAI's full-time Applications-CEO seat to a part-time adviser role (health), splitting product/ads/apps across three execs — a reason not to bank on aggressive OpenAI ad-product timelines this half. Backdrop: **SK Hynix's Nasdaq debut** (~7× oversubscribed). The move: A/B every new claimant behind ONE gateway and wait for the independent leaderboard before you reprice — don't take 'beats Opus' on a founder's word.
- **Sharpen your edge:** the connective thread to last week ('read the gauge, feed one clean input') runs straight into this one — the vendors upgraded the gauge into a *pitch*, and the operator job is to refuse the pitch on faith and run the one cheap proof. corner-shape ships because the old rendering is the fallback, so you lead with the graceful-degrade primitive and stage the flag-only one behind a guard. The CRO answer is that discipline pointed at money: don't guess the 'one thing,' let AI read YOUR order distribution, then pull one lever and watch margin, not just AOV. GPT-Live turns a walk into a thinking session — but the human still verifies the fact it fetched and keeps the confidential brief off it. And the landscape rhymes: a credibly-Opus model at a third of the price plus a second big-tech coding API mean 'which model' is a routing decision you A/B behind one gateway — not a thing you hard-code on a tweet. Take the upgrade. Run the proof. Then trust it.

The week in one idea: take the upgrade, run the one cheap proof

The vendors turned the gauge into a pitch — so the edge is the one cheap test that proves the claim on your own ground

Line the week up and it rhymes. Elon Musk called Grok 4.5 'roughly comparable to Opus 4.7 but much faster' at a third of the price — a claim, not a benchmark you ran. OpenAI's GPT-Live answers out loud from a live web search — facts it fetched, not facts you checked. CSS border - shape is 'here' — in a Canary flag, not in a browser your users have. And a teammate asked for 'the best one thing for CRO' — a tactic someone else's store proved, not yours. Last week the platforms cracked open and handed you a gauge to **read** (SNDS, PMax Channel Diagnostics, your own eval on Sonnet 5). This week they handed you a louder *claim*, and the operator move is the same discipline one notch upstream: don't reprice, redraw or re-platform on the vendor's word — run the single cheap test that turns the claim into evidence on your own ground.

Synthesis, not recap: the web-design, CRO, AI-tooling and news beats this week are four versions of one refusal — to trust a launch benchmark, a founder's tweet, a fetched fact, or a case-study 'winner' before it survives contact with your own data, your own eval, your own deploy tier, or your own verify gate.

And the test is always cheap — an eval run, a caniuise check, a 90-day export, a second source

The twin of refusing the claim is how little it costs to check it. 'Opus-class at a third of the price' becomes real or not the moment you A/B Grok 4.5 behind your gateway on three of your own tasks — and it's free to wait for an independent leaderboard until then. 'border-shape is here' becomes shippable or not the moment you read the deploy tier (flag-only → guard it; graceful-degrade → ship it). 'The one thing for CRO' becomes YOUR one thing the moment AI reads your own order distribution and names your biggest dollar leak. And GPT-Live's live answer becomes trustworthy the moment you check the source it pulled. None of these is a research project; each is a ~30-minute proof that converts a vendor's confidence into your own. Take the upgrade — the tools genuinely got better this week — but let the cheap test, not the pitch, decide when you trust it.

Do this: This week's homework, one line each: A/B **Grok 4.5** (and Muse Spark 1.1) behind your gateway on 3–5 real tasks before you reprice anything — and wait for a third-party leaderboard, not the launch tweet (AI); ship **corner-shape** today but read the Baseline/deploy tier before you lean on **border-shape**, guarding the flag-only half

behind @support's (design); hand AI your **90-day order export** + GA4 funnel and let it name your single biggest dollar leak, then pull ONE lever and watch margin (CRO); use **GPT-Live** as a brainstorm partner but verify every fact it fetches and keep confidential briefs off it (AI/all).

What changed, and what to do — by domain

One tight synthesis per domain that moved this week. The throughline (take the upgrade, run the one cheap proof) shows up differently in each.

Web design — non-rectangular UI with no SVG or clip-path: ship the half that degrades, guard the half that doesn't

[web-design · css · corner-shape · border-shape · progressive-enhancement]

Two CSS properties remove the oldest tax in UI: anything other than a rounded rectangle — a scalloped ticket, a notched card, a beveled tab, a speech-bubble — used to cost an inline SVG, a PNG mask, or a clip-path hack (which masks the box but **kills the border and box-shadow**). **corner-shape** (Chrome/Edge/Opera 139+) is the one you ship today: alongside a normal border-radius it reshapes each corner into **bevel**, **notch**, **scoop**, **squirrel** or an arbitrary superellipse(k), and takes up to four values for per-corner control (corner-shape: scoop scoop round round = a ticket). It **degrades gracefully** — any browser without it renders your existing border-radius, so there's no @support's gymnastics and no breakage, pure upside. The near-future property, **border-shape**, is more powerful: instead of masking it **redefines the box geometry**, so the **border**, **box-shadow**, **outline** and **pointer events all follow the new silhouette** (real speech-bubble tails, wavy banners with a live border). But it's a W3C Working Draft, on **Interop 2026**, and **flag-only in Chrome Canary 146+** today — nothing in stable, nothing in Firefox or Safari, syntax may still change. The gauge is the deploy tier: graceful-degrade primitive → ship now; flag-only primitive → stage behind a guard.

Do this: Add corner-shape next to a radius you already use (`.btn { border-radius: 16px; corner-shape: bevel; }`) — no guard needed; cut ticket/tab silhouettes with per-corner values (scoop scoop round round) and **notch** for coupon cut-outs. Build border-shape components now but ship them only behind @support's (border-shape: circle()) as progressive enhancement, keeping a base border-radius everyone gets and clip-path as your cross-browser fallback for hard shape needs. Preview border-shape in Chrome Canary 146+ (experimental web-platform features on) — expect bugs and syntax churn.

Keeper distinction: corner-shape is the graceful-degrade branch (ship day-one, Chromium-only, NOT Baseline); border-shape is the flag-only branch (build now, don't deploy). Logged a prediction (border-shape does NOT reach Baseline by end of 2026). Card: card-webdesign-border-shape (thread-modern-css-primitives). Cross-checked across CSS-Tricks, Una Kravets (una.im), modern-css.com, MDN, Smashing Magazine.

CRO & AOV — there's no universal 'one thing'; let your own order data name the biggest dollar leak, then pull ONE lever

[cro · aov · analytics · checkout · free-shipping]

A standing team Ask — 'what's the best ONE thing for CRO and AOV?' — got the honest answer: there's no universal one, there's the one thing *your* numbers say leaks most, and AI makes that a sub-hour diagnosis. Export your last-90-day order totals (Shopify **Admin** → **Orders** → **Export**) + your GA4 checkout funnel, hand both to Claude/ChatGPT, and ask it to compute your **AOV** and your **modal** (most-common) order value, show the distribution, and say whether you lose more money to **small baskets** or **checkout drop-off**. That single question replaces the debate with a dollar-weighted worklist. If the leak is checkout drop-off, the evidence-backed fix is killing cost-shock and forced friction: Baymard's 50-study dataset puts average cart abandonment at **~70%**, and the top fixable reasons are **unexpected extra costs (39%)**, **forced account creation (19%)**, a **too-long checkout (18%)** and **no total shown up front (14%)** — so show landed cost early, offer guest checkout, cut fields. If the leak is small baskets, peg a free-shipping threshold **just above your modal order value** (not a round number, not the blunt '30% above AOV' rule) with a live 'You're \$X away' progress bar (goal-gradient). This is 'prove it on your own ground' aimed at money: your distribution picks the lever, not someone else's case study.

Do this: Pull the data, have AI compute AOV + modal order value and rank the leak, ship ONE lever pegged to the modal value, and measure AOV, conversion AND margin together for 3–4 weeks — a higher AOV you fund with absorbed shipping can be a net loss. Card: card-cro-ai-diagnose-one-lever (updated Jul 10 to the current Baymard figures).

Answers an open Ask-anything gap (mark addressed). Anchor stats verified against the current Baymard Cart-Abandonment page (70.22%; 39/19/18/14) — corrected a stale 48%/26% web snippet; Shopify backs the peg-to-modal-value threshold. Directional: progress-bar/upsell lift %s are vendor self-report — the mechanism (goal-gradient) is sound, the number is yours to test. Card: card-cro-ai-diagnose-one-lever.

AI tooling — voice finally got current knowledge; use it to think out loud, but keep the human verify gate on what it fetches

[ai-tooling · openai · gpt-live · voice · model-router]

OpenAI shipped **GPT-Live** (Jul 8), replacing ChatGPT's old Advanced Voice Mode (a ~2024-cutoff model that made real thinking useless). It's **full-duplex** — it listens and speaks at the same time, so you interrupt it like a colleague and it stops clean, or pause mid-sentence and it waits — it **backchannels** ('mhm'), and for anything needing fresh facts or reasoning it **silently hands off to GPT-5.5 in the background** and keeps talking while the answer comes back. So voice finally has current knowledge + live web search + memory: a hands-free brainstorm or research-interview partner you can walk-and-think with. Run it today by updating the ChatGPT app (or chatgpt.com) and tapping the Voice button — no setting to flip; Free gets GPT-Live-1 mini, any paid tier gets GPT-Live-1 (iOS/Android/web/CarPlay). The honest scope is the guardrail: **no screen-share, video or vision at launch** (that's still legacy voice on the old model), **no API yet**, and it's **excluded from Business/Enterprise/Edu** — so this is a personal-productivity tool, not a client-data tool. Same 'prove the fetched fact' discipline as the model beats: the machine now speaks with confidence and current data, and the operator's job is to check the source before the claim leaves the room.

Do this: Prime it once ('be my brainstorm partner; look things up and keep talking'), then think out loud on a real brief and have it pull live competitor references without breaking rhythm — but verify every fact it fetched against a real source, and keep confidential client material off it (no ZDR/enterprise path yet). Card: card-ai-tooling-gpt-live (thread-ai-voice-agents).

Confirmed ≥4 (OpenAI + Simon Willison + TechCrunch + TestingCatalog + 9to5Mac). Logged a prediction (GPT-Live gains video/screen-share AND/OR a realtime API by ~Sep 30). Reinforces the durable lesson: a tool is only as good as its knowledge freshness + the human verify gate.

News — the model price war opened a new front and OpenAI's product org lost its owner; 'which model' is a routing decision, not a hard-code

[news · markets · grok · openai · model-router]

Two shifts worth accounting for. First, the coding-model price war got a new front. **xAI shipped Grok 4.5** (Jul 8), an 'Opus-class' model Musk pegs as 'roughly comparable to Opus 4.7 but much faster' at **\$2/\$6** per M tokens vs Opus 4.7's \$5/\$25 — the first major model since xAI's **\$60B all-stock Cursor (Anysphere) buy** (the 'trained on Cursor data' + exact SWE-Bench claims are single-source → directional; EU availability lags). A day later **Meta entered with Muse Spark 1.1**, a low-cost agentic model on a new **OpenAI/Anthropic-compatible Meta Model API** (Zuckerberg's first X post in ~3 years). And **SambaNova** closed a **\$1B first tranche at an \$11B valuation** (Series F, General Atlantic, 4-source confirmed) — capital still flooding inference silicon that isn't Nvidia. Second, **Fidji Simo is stepping down** from OpenAI's full-time Applications-CEO role to a part-time adviser seat (worsening health), splitting product, ads and apps across Brockman, Friar and Kwon — a reason not to bank on aggressive OpenAI ad-product timelines this half. Backdrop: **SK Hynix's Nasdaq debut** (when-issued 'SKHYV' Fri, regular 'SKHY' Mon; ~7× oversubscribed) — capital still chasing AI memory. The durable read: with a credibly-Opus model at a third of the price and a second big-tech coding API, 'which model' is now a routing decision you A/B behind one gateway — never repriced on a founder's tweet.

Do this: Treat every new 'beats Opus/GPT' model as a candidate, not a switch: A/B it behind your gateway (one key, every model) on 3–5 real tasks and wait for an independent leaderboard (SWE-Bench Verified, LMArena) before you reprice; build client AI for **portability** so a cheaper credible model is a config change; and lead pitches with cost-per-outcome + a spend ceiling. Card: card-ai-tooling-model-portability.

Markets figures directional (CNBC/Bloomberg 403'd — no same-day close asserted); load-bearing verified fact = the SK Hynix Nasdaq debut. Open Jun-8 'dip not regime change' call [due ~Jul 20] held open (SK Hynix strength supportive, not resolved); Jun-29 ≥\$10B-DEBT call [due ~Aug 24] no qualifying event (SambaNova/SK Hynix are equity). Logged a prediction (a third-party leaderboard lands Grok 4.5 OR Muse Spark 1.1 top-3 by ~Aug 7).

Techniques & tools worth adopting this week — the keepers

The survivors after killing 95%. Each is a run-it-today play with a real link — and each comes with the cheap test that earns your trust.

The shortlist

[technique]

Keeper	What it buys you	Start here + the discipline
corner-shape (Chrome/Edge/Opera 139+)	Beveled / notched / scooped / squircle UI with no SVG or clip-path	Add it next to border-radius; ships now because border-radius is the fallback (no @supports)
border-shape (Canary 146+, flag-only)	Border / shadow / outline follow a fully custom edge	Build components now; ship ONLY behind @supports; keep clip-path as the fallback
AI CRO diagnosis + modal-value threshold	Your own data names the biggest \$ leak	Export 90-day orders + GA4; compute AOV + modal value; ship one lever; watch margin
GPT-Live (OpenAI voice)	Full-duplex voice with live search + GPT-5.5 in the background	Prime it as a brainstorm partner; verify fetched facts; keep confidential material off it
Grok 4.5 / Muse Spark 1.1 (the router move)	Opus-class coding claims at a third of the price	A/B behind one gateway on real tasks; wait for a third-party leaderboard before repricing

Predictions ledger — resolved & newly on the board

Resolved this week

- No prediction reached its resolution window this week — the near-term calls all resolve NEXT week (see 'On the clock' below). We resolve at the window on closes, not on intraday prints, so nothing was force-closed. Honest ledger discipline: an open call under supportive pressure is still open.

Newly on the board (logged this week, still open)

- **[web-design]** CSS border - shape does NOT reach Baseline Newly Available by end of 2026 — it stays Chromium-flag-or-stable-only with no shipped Firefox or Safari implementation (basis: W3C Working Draft, flag-only in Chrome Canary 146+ in Jul 2026; even its more mature sibling corner - shape is still Chromium-only). Check via caniuse/MDN Baseline on the date — due ~Dec 31 2026.
- **[ai-tooling]** OpenAI adds live video / screen-sharing to GPT-Live (not just the legacy voice mode) AND/OR opens a GPT-Live realtime API — both were described as 'in development'/'planned soon' at the Jul 8 launch — due ~Sep 30 2026.
- **[news]** A public, independent coding/agentive leaderboard result (SWE-Bench Verified, LMArena coding, or similar third-party — not a vendor self-report) validates the 'beats OpenAI/Anthropic' framing by landing Grok 4.5 OR Meta Muse Spark 1.1 in the top 3 — due ~Aug 7 2026.

On the clock — resolving next week

- **■ [news] The AI-equity sell-off is a dip, not a regime change** (made Jun 8) — the Jun-8 basket (SoftBank, Samsung, SK Hynix) closes higher than its Jun 8 level by ~Jul 20. SK Hynix's ~7x-oversubscribed Nasdaq debut this week is supportive, not resolving; we judge it at the window, on closes not intraday prints.
- **■ [ai-tooling] GLM-5.2 reaches OpenRouter's coding/agentive top 3 by token volume** (made Jun 19) — confirms open-weights routing went mainstream — due ~Jul 19.
- **■ [news] Apple's 'sell the news' reverses** (made Jun 9) — AAPL closes above its Jun 8 (pre-WWDC) level as the Siri-AI narrative is digested — due ~Jul 21.

This week's action list

Concrete and assignable, per discipline. Each pairs taking the upgrade with the one cheap test that earns your trust.

- AI/all — A/B Grok 4.5 (and Muse Spark 1.1) behind your gateway on 3–5 real tasks before repricing anything; wait for a third-party leaderboard, not the launch tweet.
- AI/all — Use GPT-Live as a hands-free brainstorm partner; verify every fact it fetches and keep confidential client briefs off it (no enterprise/ZDR path yet).
- Design — Ship corner-shape today (no @supports needed); build border-shape components but deploy them only behind @supports, keeping clip-path as the cross-browser fallback.
- CRO — Hand AI your 90-day order export + GA4 funnel; compute AOV + modal order value; ship ONE lever pegged to the modal value and measure AOV, conversion AND margin for 3–4 weeks.
- News/all — Audit one client AI feature for model portability (gateway + fallback) so a cheaper credible model is a config change; lead the next pitch with cost-per-outcome + a spend ceiling.

Sources

CSS-Tricks — Get Ready For the Powerful CSS border-shape Property (Jul 7 2026) · Una Kravets — border-shape: the future of the non-rectangular web · modern-css.com — CSS border-shape (flag-only status, syntax) · MDN — corner-shape (values, superellipse, Baseline status) · Smashing Magazine — Beyond border-radius: what corner-shape unlocks · Baymard Institute — Cart Abandonment Rate (70.22%; 39%/19%/18%/14% reasons) · Shopify — Average Order Value (peg free-shipping threshold to modal order value) · Shopify — Free Shipping strategies (~80% will meet a threshold; dynamic progress bar) · OpenAI — Introducing GPT-Live · Simon Willison — GPT-Live hands-on · TechCrunch — OpenAI's new voice models (GPT-Live) · TestingCatalog — GPT-Live plan gating + no video/screen-share at launch · 9to5Mac — GPT-Live access + CarPlay · TechCrunch — SpaceXAI releases Grok 4.5 (\$2/\$6 vs \$5/\$25) · CNBC — SpaceX/xAI's \$60B Cursor (Anysphere) acquisition · Fortune — Meta Muse Spark 1.1 · TechCrunch — Meta enters the AI coding battle with Muse Spark 1.1 · TechCrunch — SambaNova \$1B at \$11B (Series F) · General Atlantic — SambaNova \$1B first close at \$11B · PYMNTS — Fidji Simo steps down to part-time adviser · Yahoo Finance — SK Hynix Nasdaq debut (~7× oversubscribed)

Weekly Briefing — Issue #7 · Week of Jul 6 – 12, 2026 · Compiled from multiple cross-checked sources. Figures for fast-moving deals are directional.